

# Ceng334

## Homework 3 Recitation

Fırat Ağış  
firat@ceng.metu.edu.tr

May, 2024



# Outline

- 1 Ext2 File system
  - Overview
  - Super Block
  - Block Group Descriptors
  - Bitmaps
  - Inodes
  - Directories
- 2 Assignment
  - Guarantees
  - Objectives
    - Bitmap Recovery
    - Pointer Recovery
    - User Conformation
- 3 Final Remarks



# Overview





Figure: Overall ext2 Layout

# Overview

## Boot Block




- Not relevant for our purposes.



# Overview

## Super Block




- Properties of the whole file system (ie. block size).
- Generic properties of block groups (ie. block per block group)



# Overview

## Block Group Descriptors




- Locations of important blocks in that block group.
- Properties specific to the current state of the block group.



# Overview

## Bitmaps




- One for blocks, one for inodes
- Records whether a block or inode is in use or not.
- Specific to that block group.



# Overview

## Inode Table




- Houses inodes located in the block group.
  - There are inodes for each file and directory.
  - Inodes records the properties of file/directories, as well as where their data is located within the filesystem





# Overview

## Data Blocks



- Blocks that contains data or pointers.



# Super Block

This is not the complete picture, just the relevant fields.

```
struct ext2_super_block {  
    uint32_t inode_count; /* Total number of inodes in the fs */  
    uint32_t block_count; /* Total number of blocks in the fs */  
    ...  
    uint32_t free_block_count; /* Number of free blocks */  
    uint32_t free_inode_count; /* Number of free inodes */  
    uint32_t first_data_block; /* The first data block number */  
    ...  
}
```

- The super block always begins at byte 1024.
- `inode_count` and `block_count`: Total number of blocks and inodes in the file system.
- `free_block_count` and `free_inode_count`: How many of those inodes and blocks are not in use.
- `first_data_block`: Number of the first data block, depends on the block size.



# Super Block

```
uint32_t log_block_size; /* 2^(10 + this value) gives the block size */  
...  
uint32_t blocks_per_group; /* Number of blocks for each block group */  
...  
uint32_t inodes_per_group; /* Number of inodes for each block group */  
...  
uint32_t first_inode; /* First non-reserved inode  
                      in the file system */  
uint16_t inode_size; /* Size of each inode */  
/* More, less relevant fields follow */  
};
```

- `log_block_size`: Size of each block, given by the formula.
- `blocks_per_group` and `inodes_per_group`: How many blocks and inodes are within a block group. Last block group can have less than these.
- `first_inode`: First inode not reserved by the system.
- `inode_size`: Size of each inode.




# Block Group Descriptors

```
struct ext2_block_group_descriptor {  
    uint32_t block_bitmap; /* Block containing the block bitmap */  
    uint32_t inode_bitmap; /* Block containing the inode bitmap */  
    uint32_t inode_table; /* First block of the inode table */  
    uint16_t free_block_count; /* Number of free blocks in the group */  
    uint16_t free_inode_count; /* Number of free inodes in the group */  
    uint16_t used_dirs_count; /* Number of directories in the group */  
    uint16_t pad; /* Padding to 4 byte alignment */  
    uint32_t reserved[3]; /* Unused, reserved 12 bytes */  
};
```

- Always starts at the first unoccupied block after the super block.
- You should read the locations of bitmaps and inode tables from here.



# Bitmaps



- Bitmaps are always exactly 1 block and there are only 1 bitmap of each kind in a block group.
- A bit corresponding to an inode or a block is set to 1 if its in use, set to 0 if not.
- Each bitmap only contains information about the blocks/inodes in the block group it's in.



# Inodes

```
struct ext2_inode {  
    uint16_t mode; /* Contains filetype and permissions */  
    uint16_t uid; /* Owning user id */  
    uint32_t size; /* Least significant 32-bits of file size in rev. 1 */  
    uint32_t access_time; /* Timestamps (in seconds since 1 Jan 1970) */  
    uint32_t creation_time;  
    uint32_t modification_time;  
    uint32_t deletion_time; /* Zero for non-deleted inodes! */  
    ...  
};
```

- mode: File type (user file or directory for our purposes) and permissions for the file.
- size: Least significant 32-bits of file size.
- access\_time, creation\_time, and modification\_time: Various timestamps.
- deletion\_time: Also a timestamp. Zero if an inode is not deleted.



# Inodes


```
...
uint16_t gid; /* Owning group id */
uint16_t link_count; /* Number of hard links */
uint32_t block_count_512; /* Number of 512-byte blocks
alloc'd to file */
uint32_t flags; /* Special flags */
uint32_t reserved; /* 4 reserved bytes */
...
};
```

- uid and gid: User id and Group id of the user and the group that owns the file.
- block\_count\_512: Number of blocks allocated to the file if block size was 512-bytes. Can be use to calculate how many blocks are allocated to a file.



# Inodes

```
...
uint32_t direct_blocks[12]; /* Direct data blocks */
uint32_t single_indirect; /* Single indirect block */
uint32_t double_indirect; /* Double indirect block */
uint32_t triple_indirect; /* Triple indirect block */
/* Other, less relevant fields follow */
};
```





# Directories

```
struct ext2_dir_entry {  
    uint32_t inode; /* inode number of the file */  
    uint16_t length; /* Record length, aligned on 4 bytes */  
    uint8_t name_length; /* 255 is the maximum allowed name length */  
    uint8_t file_type; /* Not used in rev. 0,  
                        file type identifier in rev. 1 */  
    char name[]; /* File name. This is called a  
                 'flexible array member' in C. */  
};
```

- Each directory contains the entries for "." and "..".
- If inode field is 0, the directory entry should be ignored.
- length field of the last directory entry in a block is set in a way that it fills the remaining space in the block.



# Assignment

- You will be designing a recovery program for ext2 file system.
- You will be using C or C++ programming languages.
- Your assignment will be tested on inek machines.
- Your due date is 30 May 2024, Thursday, 23:59

# Guarantees

- The data erasure is limited to bitmaps and pointer fields of inodes, everything else is completely intact.
- Contents of the inodes reserved by the system is also intact.
- Any unused data block is wiped clean, and contains no set bits.
- Any non-empty data block that contains user data (ie. not directory or pointer data) starts with the same 32-byte identifier that will be given at runtime.
- Pointers to data blocks that are empty are intact.
- First two parts of the homework are independent, but completing one might help you with the other one.



# Objectives

Your goal is:

- Recover deleted bitmap bits.
- Recover deleted pointers on inodes.
- Print the structure of the file system.



# Objectives

## Bitmap Recovery

Bitmaps on the file system might be damaged.

- That damage can be partial.
- Not all bitmaps might not be damaged.
- Bitmap damage is limited to bits being set to 0.



# Objectives

## Pointer Recovery

Some pointers of inodes are set to 0 incorrectly, you need to fix it using (not limited to) the following informations:

- Data blocks in use that are not being pointed to.
- Inodes of "." and ".." directory entities.
- Sizes of files located in inodes.



# Objectives

## User Conformation

Finally you will print the recovered file system to the standard output in a tree structure, such as:

```
- root/  
-- home/  
--- user1/  
---- file1.txt  
---- file2.txt  
--- user2/  
---- file3.txt  
-- etc/  
--- config.txt
```



# Final Remarks

- This presentation is to aid you, but it is not comprehensive, refer to the homework text for any specifics.
- For simplification, the outputs can be in any order.
- Including POSIX ext2/ext3/ext4 libraries (as well as kernel codes and their variations etc.) is **not** allowed.
- There will be no extension for this homework, use your time wisely.

Good Luck Everyone